FINSHIELD AI

Fraud Risk Decision Policy

Controls for explainable hybrid transaction risk decisions

Version	1.0
Classification	Internal demonstration
Owner	Fraud Operations

Policy principle: a high score prioritizes investigation; it is not a substitute for verified evidence.

Purpose and scope

Objective

Prioritize suspicious digital payments using model probability, deterministic policy rules, behavioral signals, and investigator review. The system supports decisions; it does not establish that fraud occurred.

Data boundary

PaySim is synthetic. Production use requires representative data, calibration, monitoring, access controls, and governance approval.

Decision framework

Unified score

Default weighting is 65 percent ML, 20 percent rules, and 15 percent behavioral risk. Scores 0-30 approve, above 30-60 monitor, above 60-80 manual review, and above 80 block.

Evidence standard

Every decision must retain model version, probability, threshold, feature snapshot, triggered rules, SHAP factors, behavioral signals, final score, and action.

Overrides

Investigators may override an automated action only with a documented reason and supporting evidence. Overrides must remain auditable.

Risk controls

False positives

Review trusted beneficiaries, known customer behavior, balance corrections, and successful step-up verification before maintaining restrictions.

False negatives

Monitor missed-fraud rate, captured fraud value, emerging typologies, threshold performance, and drift by transaction type.

Explainability

SHAP describes model behavior and is not causal proof. Correlated balance features can redistribute attribution.

Governance

Review cadence

Fraud operations should review thresholds, rule hit rates, model calibration, investigation capacity, and customer friction on an approved schedule.

Prohibited use

Do not use PaySim labels, synthetic relationships, or this demonstration policy for autonomous real-world adverse action.